

DBMS_XPLAN cheat sheet

Method #1: EXPLAIN PLAN

`explain plan [set statement_id = 'xxx'] for SELECT_or_DML_query;`

This computes the plan of the target query (without running it), and inserts the details of the plan into the PLAN_TABLE; then:

```
select * from table(dbms_xplan.display('PLAN_TABLE', 'xxx', 'display_fmt'));
```

Or simply, if not using a statement identifier:

```
select * from table(dbms_xplan.display(format => 'display_fmt'));
```

Requirements

- Permissions to run the target SELECT or DML statement
- If using views: READ or SELECT privilege on underlying tables (otherwise ORA-01039 is raised)

Display formats

- Default: typical; but use at least: +alias (= 'typical +alias')
- Preferred: all -projection (this adds the “Hint Report” section if DB ≥ 19c)
- Append: +outline, or use: advanced -projection to show the outline data; if using the latter, also add: -qbregrity if DB ≥ 19c, to turn off the Query Block Registry section
- Append: +projection if you need the Column Projection section
- Add: +adaptive to fully show adaptive plans, if the DB is configured to use them
- DB ≥ 23c: the new “SQL Analysis Report” section may be turned off if needed, by appending -sql_analysis_report

Limitations

- explain plan does not use *bind peeking*, therefore it may generate plans which differ from actual plans if *bind variables* are being used, and column histograms exist
- DBAs beware: the MERGE ANY VIEW privilege (included in the DBA role) causes *secure view merging* to be bypassed, creating opportunities for explain plan to generate plans unavailable to users subjected to *secure view merging*
- explain plan always uses the latest (published) statistics—possibly causing it to not reproduce plans created earlier in the cursor cache

Notes

- The (estimated) Bytes, Rows, Cost, and Time columns are *per-execution* of the concerned plan line operation (on the contrary, the columns Starts, A-Rows, A-Time, Buffers, and Reads, are *cumulative*)
- Caution: explain plan *inserts* rows into the PLAN_TABLE, which means that if you didn't have an on-going transaction *before* explain plan, you have one *after*.

Method #2: DBMS_XPLAN.DISPLAY_CURSOR

`SELECT_or_DML_query;`

... then retrieve the *sql_id* + *child_number* of the cursor, and then:

```
select * from table(dbms_xplan.display_cursor('sql_id', child_number, 'display_fmt'));
```

Special case: latest cursor in this session ^(†) ^(‡)

```
select * from table(dbms_xplan.display_cursor(null, null, 'display_fmt'));
```

Requirements

- READ / SELECT privilege on V\$SQL, V\$SQL_PLAN_STATISTICS_ALL
- ^(†) Additional privilege required: READ / SELECT on V\$SESSION
- ^(‡) set serveroutput off is needed for this to work as expected

Display formats

- Same formats as in dbms_xplan.display
- Use: +peeked_binds (or advanced) to view the values, if any, of binds peeked at parse time
- Add: +iostats | memstats | allstats [all | last] to print *actual* execution statistics:
 - all | last: show statistics for all executions / for the last execution (default: all)
Remark: statistics from *on-going* executions are not available
 - +memstats: adds the 0Mem, 1Mem, Used-Mem (or 0/1/M) columns
 - +iostats: adds the Starts, A-Rows, A-Time, Buffers, Reads, and Writes columns
 - +allstats = both iostats and memstats

Note: +iostats requires either:

- alter session set statistics_level = all; caution: high overhead due to timing code, making queries possibly slower—for tests only!

Or:

- the /*+ gather_plan_statistics */ hint, which incurs lower timing overhead, but can result in the A-Time column being (very) inaccurate

Limitations

- The security requirements make this a privileged operation: users with the required permissions may see every plan (and every query) database wide
- In some cases, dbms_xplan.display_cursor does not correctly render complex predicates, yielding meaningless expressions such as: filter(IS NULL), access("COL"=PRIOR NULL), etc. This is a limitation in V\$SQL_PLAN_STATISTICS_ALL, which cannot correctly render such predicates—only explain plan can
- +allstats last does not function correctly with parallel queries: execution statistics from PX processes are not added to those of the QC; workaround: use +allstats (all implicitly)

Reminder: always pay attention to the Notes section, if any

Plan table columns

Column Name	Description
Id	Identifying number of the operation in the plan; in the case of adaptive plans, use +adaptive to display the “true” plan line identifiers from the full (unresolved) plan
Operation	Name of the operation, along with its execution options
Name	Name of the object that is read from / written into
Starts	Actual number of times that this operation was started; special case: number of times that the memory structure of the 2nd child of a MERGE JOIN was probed
Rows / E-Rows	Estimated number of rows to be produced (note: <i>per execution</i>)
Bytes / E-Bytes	Estimated number of bytes to be produced (note: <i>per execution</i>)
TempSpc / E-Temp	Estimated amount of temp. storage required for this operation, in bytes (note: <i>per execution</i>)
Cost (%CPU)	Estimated “cost” of the operation & percentage of CPU time thereof, if available (note: <i>per execution</i>)
Time / E-Time	Estimated duration of the operation, in hours:minutes:seconds (note: <i>per execution</i>); remark: values are usually inaccurate if system statistics have not been properly configured
Pstart / Pstop	Partition number of the 1st / last partition iterated upon in this operation, if pre-determined; otherwise the method of determination to be used at runtime, e.g.: KEY, KEY(I), KEY(SQ), :BFnnnn, etc.
TQ	In parallel queries, identifier of the DFO (<i>Data Flow Operation</i>) that this operation belongs to; it is also the name of the “Table Queue” that PX servers running this operation write into
IN-OUT	In parallel queries, relationship between this operation and related operations
PQ Distrib	In parallel queries, method of distributing rows in inter-group communications
A-Rows	Actual number of rows produced by this operation (<i>cumulative</i>) Note: for operations on bitmap indexes, this is <i>not</i> the count of table rows!
A-Time	Actual time spent running this operation and its child operations (<i>cumulative</i>) Keep in mind that statistics_level = all causes distortion due to time-measurement code overhead, whereas the alternative, the /*+ gather_plan_statistics */ hint, may return false durations
Buffers	Actual count of logical blocks read by this operation, either from the buffer cache, or from datafiles if doing direct path reads (<i>cumulative</i>)
Reads	Actual count of blocks physically read, if any, by this operation (<i>cumulative</i>)
Writes	Actual count of blocks written, if any, by this operation (<i>cumulative</i>)
OMem	Work area size required for running this operation fully in memory; unit: bytes
1Mem	Work area size required for running this operation using a one-pass (memory + temp. space) algorithm; unit: bytes
Used-Mem	Actual amount of memory used in the last execution of this op.; unit: bytes
0/1/M	Number of times that this operation was run, respectively, fully in memory, using a one-pass algorithm, or using a multi-pass algorithm (<i>cumulative</i>)
Used-Tmp	Amount of temporary storage used when performing this operation; unit: bytes (<i>cumulative</i>); note: the unit is <i>kilobytes</i> in DB versions ≤ 12.2
Max-Tmp	Maximum amount of temp. storage used when performing this operation; unit: bytes (<i>cumulative</i>); note: the unit is <i>kilobytes</i> in DB versions ≤ 12.2

Plan table columns (continued)

Column Name	Description
Inst	In distributed queries, name of the DB link used by this operation

Parallel Queries

Input / Output

IN-OUT	Description
P->S	Parallel to Serial: this parallel operation sends rows to the QC (<i>query coordinator</i>) of its DFO tree
S->P	Serial to Parallel: this serial operation sends rows to a parallel operation (remark: this could be inefficient, due to this serial operation possibly acting as a bottleneck)
P->P	Parallel to Parallel: this parallel operation sends rows to another parallel operation
PCWC	Parallel combined with Child: each PX process in this operation receives rows from its child operation in the same process (no inter-process communication happens)
PCWP	Parallel combined with Parent: each PX process in this operation sends rows to its parent operation in the same process (no inter-process communication happens)
SCWC	Serial combined with Child: this serial operation runs in the same process as its child operation, which it receives rows from (no inter-process communication happens)
SCWP	Serial combined with Parent: this serial operation runs in the same process as its parent operation, which it returns rows to (no inter-process communication happens)

Distribution Methods

PQ Distrib	Description
QC (RAND)	The QC (query coordinator) is being sent rows with no ordering
QC (ORDER)	The QC (query coordinator) consumes rows in order, beginning with rows from the first PX process, and ending with rows from the last PX process
BROADCAST	Every producer sends all its rows to every consumer
BROADCAST LOCAL	RAC-only: same as Broadcast, but producers send to consumers in the local instance
RANGE	Every producer sends each row to a <i>single</i> consumer, based on a range condition: rows in the same range are sent to the same consumer
RND-ROBIN	Each producer sends each row to one consumer in turn (resulting in rows being evenly distributed among consumers)
HASH	Every producer sends each row to a <i>single</i> consumer, based on a hash function; rows with the same hash value are sent to the same consumer
HYBRID HASH	Producers send rows to consumers using either Broadcast or Hash, the choice between the two being determined at runtime (DB ≥ 12.1)
PART (ROWID)	Rows are sent to consumers based on the partitioning of a table/index, using the rowid of the row to update/delete
PART (KEY)	Rows are sent to consumers based on the partitioning of a table/index, using the partition key
RANDOM LOCAL	A variation of PART (KEY) in which sets of consumer processes are maintaining sets of partitions
1 SLAVE	All rows are sent to a single process in the consumer group (DB ≥ 12.1)