

HASH JOIN pseudo-code *(high-level, simplified perspective)*

HASH JOIN

CHILD_ROW_SOURCE_1	← driving/build row source, or “left” input	alias: r_1	columns: (c_1, c_2, \dots, c_n)
CHILD_ROW_SOURCE_2	← probe row source, or “right” input	alias: r_2	columns: (c_1, c_2, \dots, c_m)

with join conditions as follows:

$$\left. \begin{array}{l} r_1 \cdot c_{h_1} = r_2 \cdot c_{j_1} \\ \text{and } r_1 \cdot c_{h_2} = r_2 \cdot c_{j_2} \\ \dots \\ \text{and } r_1 \cdot c_{h_k} = r_2 \cdot c_{j_k} \end{array} \right\} \text{equality conditions}$$

$$\text{and } \left. \begin{array}{l} \text{expression}(r_1 \cdot c_{h_{k+1}}, \dots, r_1 \cdot c_{h_p} \\ \quad \quad \quad , r_2 \cdot c_{j_{k+1}}, \dots, r_2 \cdot c_{j_q}) \end{array} \right\} \text{non-equality conditions}$$

Start CHILD_ROW_SOURCE_1

```

For each row  $r_1 = (c_1, c_2, \dots, c_n)$  from CHILD_ROW_SOURCE_1 Loop -- build loop
    insert  $r_1$  into the hash table using  $(r_1.c_{h_1}, \dots, r_1.c_{h_k})$  as the hash key

```

```
End loop -- CHILD_ROW_SOURCE_1 has been fully processed
```

If CHILD_ROW_SOURCE_1 returned at least 1 row Then

```
Start CHILD_ROW_SOURCE_2
```

For each row $r_2 = (c_1, c_2, \dots, c_m)$ from CHILD_ROW_SOURCE_2 **Loop** -- probe loop

For each row r_1 matching $(r_2.c_{j_1}, \dots, r_2.c_{j_k})$ in the hash table */* access conditions */* **Loop**

```
/* evaluate non-equality conditions: filter conditions */
```

If $\text{expression}(r_1.c_{h_{k+1}}, \dots, r_1.c_{h_p}, r_2.c_{j_{k+1}}, \dots, r_2.c_{j_q})$ **is true** **Then**

Yield the combined row $\mathbf{rj} = (r_1.c_1, \dots, r_1.c_n, r_2.c_1, \dots, r_2.c_m)$ to the parent operation $(*)$

End If

End Loop

End Loop

End If

(*) Actually, only projected columns are passed to the parent operation

Key points:

- CHILD_ROW_SOURCE_1 and _2 are started only once (per start of the parent), and processed independently, in turn
- The hash table (in workarea) is built from CHILD_ROW_SOURCE_1: rows from CHILD_ROW_SOURCE_2 are not buffered (*iff* the hash join can be processed fully in memory)
- The hash key is formed of equi-joined columns; non-equality join conditions are always used as *filter* conditions, and evaluated by *iterating* on rows matching the probe key in the hash table—if there are too many such rows, a lot of CPU time could go into that
- The optimizer may swap join inputs, depending on (estimated) memory requirements of using either as the build row source